Assignment 3.1 Exercise

In this assignment, you will explore the impact of fairness metrics after applying the Disparate Impact Remover (DIR), retrain a Light GBM model with the adjusted data, and compare its metrics (listed below) against the original and reweighted models from the lab tutorial. Afterward, respond to the reflection questions.

Part 1: Experimentation with the DIR

- Modify the levels array for the Disparate Impact Remover to include finer and coarser granularity. (See example in lab/assignment notebook.)
- Identify the best repair level and the Disparate Impact at the best repair level.
- Retrain a Light GBM model on the modified data from the DIR and compute both the performance and fairness metrics.

1.1a Code Output for Finer Granularity

Answer

The output is organized as below:

- **Finer Granular and Coarser Levels array Information**
- **Combined Finer Granular and Coarse Disparate Impact Results for different levels**
- **Combined Finer Granular and Coarse Disparate Impact Plots**
- **Finer Granular Disparate Impact Plot**
- **Finer Granular Disparate Impact Results**

Finer Granular and Coarser Levels used for Disparate Impact Remover (DIR) are provided below

```
Granular Levels: [0.      0.00125 0.0025  0.00375 0.005   0.00625 0.0075  0.00875 0.01
 0.01125 0.0125  0.01375 0.015   0.01625 0.0175  0.01875 0.02    0.02125
 0.0225  0.02375 0.025   0.02625 0.0275  0.02875 0.03    0.03125 0.0325
 0.03375 0.035   0.03625 0.0375  0.03875 0.04    0.04125 0.0425  0.04375
 0.045   0.04625 0.0475  0.04875 0.05    0.05125 0.0525  0.05375 0.055
 0.05625 0.0575  0.05875 0.06    0.06125 0.0625  0.06375 0.065   0.06625
 0.0675  0.06875 0.07    0.07125 0.0725  0.07375 0.075   0.07625 0.0775
 0.07875 0.08    0.08125 0.0825  0.08375 0.085   0.08625 0.0875  0.08875
 0.09    0.09125 0.0925  0.09375 0.095   0.09625 0.0975  0.09875 0.1    ]

Coarse Levels: [0.2        0.24210526 0.28421053 0.32631579 0.36842105 0.41052632
 0.45263158 0.49473684 0.53684211 0.57894737 0.62105263 0.66315789
 0.70526316 0.74736842 0.78947368 0.83157895 0.87368421 0.91578947
 0.95789474 1.        ]
```

```
DIR at different Repair levels are provided below
```

```
  1%|          | 1/101 [00:21<36:16, 21.77s/it]Repair Level: 0.0000, Disparate Impact: 0.8350
  2%||          | 2/101 [00:37<30:01, 18.20s/it]Repair Level: 0.0013, Disparate Impact: 0.8434
  3%||          | 3/101 [00:45<21:59, 13.47s/it]Repair Level: 0.0025, Disparate Impact: 0.8464
  4%||          | 4/101 [00:54<18:46, 11.61s/it]Repair Level: 0.0037, Disparate Impact: 0.8467
  5%||          | 5/101 [01:02<16:49, 10.51s/it]Repair Level: 0.0050, Disparate Impact: 0.8401
  6%||          | 6/101 [01:10<15:05,  9.53s/it]Repair Level: 0.0063, Disparate Impact: 0.8448
  7%||          | 7/101 [01:19<14:47,  9.44s/it]Repair Level: 0.0075, Disparate Impact: 0.8405
  8%||          | 8/101 [01:28<14:09,  9.14s/it]Repair Level: 0.0088, Disparate Impact: 0.8367
  9%||          | 9/101 [01:36<13:32,  8.83s/it]Repair Level: 0.0100, Disparate Impact: 0.8419
 10%||          | 10/101 [01:44<13:20,  8.80s/it]Repair Level: 0.0112, Disparate Impact: 0.8456
 11%||          | 11/101 [01:52<12:38,  8.42s/it]Repair Level: 0.0125, Disparate Impact: 0.8330
 12%||          | 12/101 [02:01<12:43,  8.58s/it]Repair Level: 0.0138, Disparate Impact: 0.8420
 13%||          | 13/101 [02:10<12:49,  8.75s/it]Repair Level: 0.0150, Disparate Impact: 0.8332
 14%||          | 14/101 [02:18<12:12,  8.42s/it]Repair Level: 0.0163, Disparate Impact: 0.8411
 15%||          | 15/101 [02:26<12:13,  8.53s/it]Repair Level: 0.0175, Disparate Impact: 0.8338
 16%||          | 16/101 [02:35<11:55,  8.41s/it]Repair Level: 0.0187, Disparate Impact: 0.8443
 17%||          | 17/101 [02:43<11:43,  8.37s/it]Repair Level: 0.0200, Disparate Impact: 0.8365
 18%||          | 18/101 [02:52<11:56,  8.63s/it]Repair Level: 0.0213, Disparate Impact: 0.8367
 19%||          | 19/101 [03:00<11:23,  8.34s/it]Repair Level: 0.0225, Disparate Impact: 0.8393
 20%||          | 20/101 [03:08<11:22,  8.43s/it]Repair Level: 0.0238, Disparate Impact: 0.8381
 21%||          | 21/101 [03:17<11:25,  8.57s/it]Repair Level: 0.0250, Disparate Impact: 0.8398
 22%||          | 22/101 [03:25<10:58,  8.34s/it]Repair Level: 0.0262, Disparate Impact: 0.8440
 23%||          | 23/101 [03:34<11:11,  8.61s/it]Repair Level: 0.0275, Disparate Impact: 0.8317
 24%||          | 24/101 [03:43<10:53,  8.49s/it]Repair Level: 0.0288, Disparate Impact: 0.8336
 25%||          | 25/101 [03:50<10:27,  8.26s/it]Repair Level: 0.0300, Disparate Impact: 0.8422
```

```
 26%||          | 26/101 [03:59<10:35,  8.47s/it]Repair Level: 0.0312, Disparate Impact: 0.8390
 27%||          | 27/101 [04:07<10:11,  8.26s/it]Repair Level: 0.0325, Disparate Impact: 0.8403
cell output actions
 28%||          | 28/101 [04:16<10:16,  8.44s/it]Repair Level: 0.0338, Disparate Impact: 0.8367
 29%||          | 29/101 [04:25<10:20,  8.62s/it]Repair Level: 0.0350, Disparate Impact: 0.8363
 30%||          | 30/101 [04:32<09:43,  8.21s/it]Repair Level: 0.0362, Disparate Impact: 0.8459
 31%||          | 31/101 [04:41<09:56,  8.53s/it]Repair Level: 0.0375, Disparate Impact: 0.8439
 32%||          | 32/101 [04:50<09:39,  8.40s/it]Repair Level: 0.0387, Disparate Impact: 0.8439
 33%||          | 33/101 [04:58<09:30,  8.39s/it]Repair Level: 0.0400, Disparate Impact: 0.8417
 34%||          | 34/101 [05:07<09:42,  8.70s/it]Repair Level: 0.0413, Disparate Impact: 0.8429
 35%||          | 35/101 [05:15<09:08,  8.30s/it]Repair Level: 0.0425, Disparate Impact: 0.8447
 36%||          | 36/101 [05:24<09:09,  8.45s/it]Repair Level: 0.0438, Disparate Impact: 0.8400
 37%||          | 37/101 [05:33<09:12,  8.63s/it]Repair Level: 0.0450, Disparate Impact: 0.8406
 38%||          | 38/101 [05:40<08:45,  8.34s/it]Repair Level: 0.0462, Disparate Impact: 0.8307
 39%||          | 39/101 [05:49<08:52,  8.58s/it]Repair Level: 0.0475, Disparate Impact: 0.8439
 40%||          | 40/101 [05:57<08:31,  8.39s/it]Repair Level: 0.0488, Disparate Impact: 0.8368
 41%||          | 41/101 [06:06<08:22,  8.38s/it]Repair Level: 0.0500, Disparate Impact: 0.8377
 42%||          | 42/101 [06:15<08:33,  8.69s/it]Repair Level: 0.0513, Disparate Impact: 0.8338
 43%||          | 43/101 [06:23<08:13,  8.52s/it]Repair Level: 0.0525, Disparate Impact: 0.8327
 44%||          | 44/101 [06:32<08:07,  8.54s/it]Repair Level: 0.0537, Disparate Impact: 0.8366
 45%||          | 45/101 [06:40<07:59,  8.57s/it]Repair Level: 0.0550, Disparate Impact: 0.8407
 46%||          | 46/101 [06:48<07:37,  8.32s/it]Repair Level: 0.0563, Disparate Impact: 0.8358
 47%||          | 47/101 [06:57<07:42,  8.56s/it]Repair Level: 0.0575, Disparate Impact: 0.8362
 48%||          | 48/101 [07:06<07:37,  8.63s/it]Repair Level: 0.0588, Disparate Impact: 0.8370
 49%||          | 49/101 [07:14<07:17,  8.41s/it]Repair Level: 0.0600, Disparate Impact: 0.8369
 50%||          | 50/101 [07:23<07:15,  8.55s/it]Repair Level: 0.0612, Disparate Impact: 0.8404
 50%||          | 51/101 [07:30<06:51,  8.23s/it]Repair Level: 0.0625, Disparate Impact: 0.8259
```

```
51%|█████     | 52/101 [07:39<06:52,  8.41s/it]Repair Level: 0.0638, Disparate Impact: 0.8374
52%|█████     | 53/101 [07:48<06:54,  8.64s/it]Repair Level: 0.0650, Disparate Impact: 0.8339
53%|█████     | 54/101 [07:56<06:37,  8.46s/it]Repair Level: 0.0663, Disparate Impact: 0.8376
54%|█████     | 55/101 [08:05<06:31,  8.51s/it]Repair Level: 0.0675, Disparate Impact: 0.8403
55%|█████     | 56/101 [08:14<06:31,  8.71s/it]Repair Level: 0.0688, Disparate Impact: 0.8370
56%|█████     | 57/101 [08:22<06:13,  8.48s/it]Repair Level: 0.0700, Disparate Impact: 0.8337
57%|█████     | 58/101 [08:32<06:16,  8.76s/it]Repair Level: 0.0713, Disparate Impact: 0.8298
58%|█████     | 59/101 [08:40<06:04,  8.68s/it]Repair Level: 0.0725, Disparate Impact: 0.8325
59%|█████     | 60/101 [08:48<05:43,  8.37s/it]Repair Level: 0.0737, Disparate Impact: 0.8372
60%|██████    | 61/101 [08:57<05:42,  8.56s/it]Repair Level: 0.0750, Disparate Impact: 0.8327
61%|██████    | 62/101 [09:04<05:24,  8.31s/it]Repair Level: 0.0762, Disparate Impact: 0.8372
62%|██████    | 63/101 [09:13<05:20,  8.45s/it]Repair Level: 0.0775, Disparate Impact: 0.8477
63%|██████    | 64/101 [09:22<05:19,  8.64s/it]Repair Level: 0.0788, Disparate Impact: 0.8378
64%|██████    | 65/101 [09:30<04:57,  8.26s/it]Repair Level: 0.0800, Disparate Impact: 0.8339
65%|██████    | 66/101 [09:39<05:04,  8.70s/it]Repair Level: 0.0813, Disparate Impact: 0.8383
66%|██████    | 67/101 [09:49<05:05,  8.98s/it]Repair Level: 0.0825, Disparate Impact: 0.8285
67%|██████    | 68/101 [09:57<04:47,  8.71s/it]Repair Level: 0.0838, Disparate Impact: 0.8376
68%|██████    | 69/101 [10:07<04:48,  9.02s/it]Repair Level: 0.0850, Disparate Impact: 0.8313
69%|██████    | 70/101 [10:15<04:32,  8.78s/it]Repair Level: 0.0863, Disparate Impact: 0.8313
70%|███████   | 71/101 [10:23<04:15,  8.53s/it]Repair Level: 0.0875, Disparate Impact: 0.8284
71%|███████   | 72/101 [10:32<04:13,  8.74s/it]Repair Level: 0.0887, Disparate Impact: 0.8357
72%|███████   | 73/101 [10:40<03:55,  8.42s/it]Repair Level: 0.0900, Disparate Impact: 0.8402
73%|███████   | 74/101 [10:49<03:49,  8.49s/it]Repair Level: 0.0912, Disparate Impact: 0.8389
74%|███████   | 75/101 [10:57<03:40,  8.48s/it]Repair Level: 0.0925, Disparate Impact: 0.8428
75%|███████   | 76/101 [11:05<03:24,  8.19s/it]Repair Level: 0.0938, Disparate Impact: 0.8341

 76%|███████   | 77/101 [11:14<03:25,  8.54s/it]Repair Level: 0.0950, Disparate Impact: 0.8385
 77%|███████   | 78/101 [11:23<03:18,  8.61s/it]Repair Level: 0.0963, Disparate Impact: 0.8407
 78%|███████   | 79/101 [11:31<03:06,  8.47s/it]Repair Level: 0.0975, Disparate Impact: 0.8421
 79%|███████   | 80/101 [11:40<03:01,  8.62s/it]Repair Level: 0.0988, Disparate Impact: 0.8347
 80%|████████  | 81/101 [11:48<02:47,  8.36s/it]Repair Level: 0.1000, Disparate Impact: 0.8421
 81%|████████  | 82/101 [11:57<02:42,  8.57s/it]Repair Level: 0.2000, Disparate Impact: 0.8495
 82%|████████  | 83/101 [12:06<02:37,  8.75s/it]Repair Level: 0.2421, Disparate Impact: 0.8494
 83%|████████  | 84/101 [12:13<02:23,  8.44s/it]Repair Level: 0.2842, Disparate Impact: 0.8433
 84%|████████  | 85/101 [12:22<02:15,  8.44s/it]Repair Level: 0.3263, Disparate Impact: 0.8539
 85%|████████  | 86/101 [12:30<02:06,  8.44s/it]Repair Level: 0.3684, Disparate Impact: 0.8566
 86%|████████  | 87/101 [12:38<01:55,  8.28s/it]Repair Level: 0.4105, Disparate Impact: 0.8597
 87%|████████  | 88/101 [12:47<01:51,  8.56s/it]Repair Level: 0.4526, Disparate Impact: 0.8605
 88%|████████  | 89/101 [12:55<01:39,  8.32s/it]Repair Level: 0.4947, Disparate Impact: 0.8659
 89%|████████  | 90/101 [13:03<01:31,  8.28s/it]Repair Level: 0.5368, Disparate Impact: 0.7930
 90%|█████████ | 91/101 [13:12<01:24,  8.49s/it]Repair Level: 0.5789, Disparate Impact: 0.7894
 91%|█████████ | 92/101 [13:20<01:13,  8.12s/it]Repair Level: 0.6211, Disparate Impact: 0.7959
 92%|█████████ | 93/101 [13:29<01:07,  8.47s/it]Repair Level: 0.6632, Disparate Impact: 0.8057
 93%|█████████ | 94/101 [13:37<00:57,  8.28s/it]Repair Level: 0.7053, Disparate Impact: 0.8161
 94%|█████████ | 95/101 [13:44<00:48,  8.02s/it]Repair Level: 0.7474, Disparate Impact: 0.8368
 95%|█████████ | 96/101 [13:52<00:39,  7.96s/it]Repair Level: 0.7895, Disparate Impact: 0.8335
 96%|█████████ | 97/101 [13:59<00:30,  7.74s/it]Repair Level: 0.8316, Disparate Impact: 0.8341
 97%|█████████ | 98/101 [14:07<00:23,  7.86s/it]Repair Level: 0.8737, Disparate Impact: 0.8257
 98%|█████████ | 99/101 [14:14<00:15,  7.51s/it]Repair Level: 0.9158, Disparate Impact: 0.8142
 99%|█████████| 100/101 [14:22<00:07,  7.51s/it]Repair Level: 0.9579, Disparate Impact: 0.8074
100%|█████████| 101/101 [14:28<00:00,  8.60s/it]Repair Level: 1.0000, Disparate Impact: 0.8123
```

Results Interpretation

The results show that small fairness adjustments (Granular Levels) have minimal impact on Disparate Impact (DI), while moderate repair levels (0.3 - 0.5) improve fairness significantly. However, excessive fairness repair (>0.5) reduces DI, suggesting diminishing returns and potential performance degradation. The optimal fairness correction lies within

moderate repair levels, balancing bias mitigation without harming accuracy.

Combined Finer Granular and Coarse Disparate Impact Plots

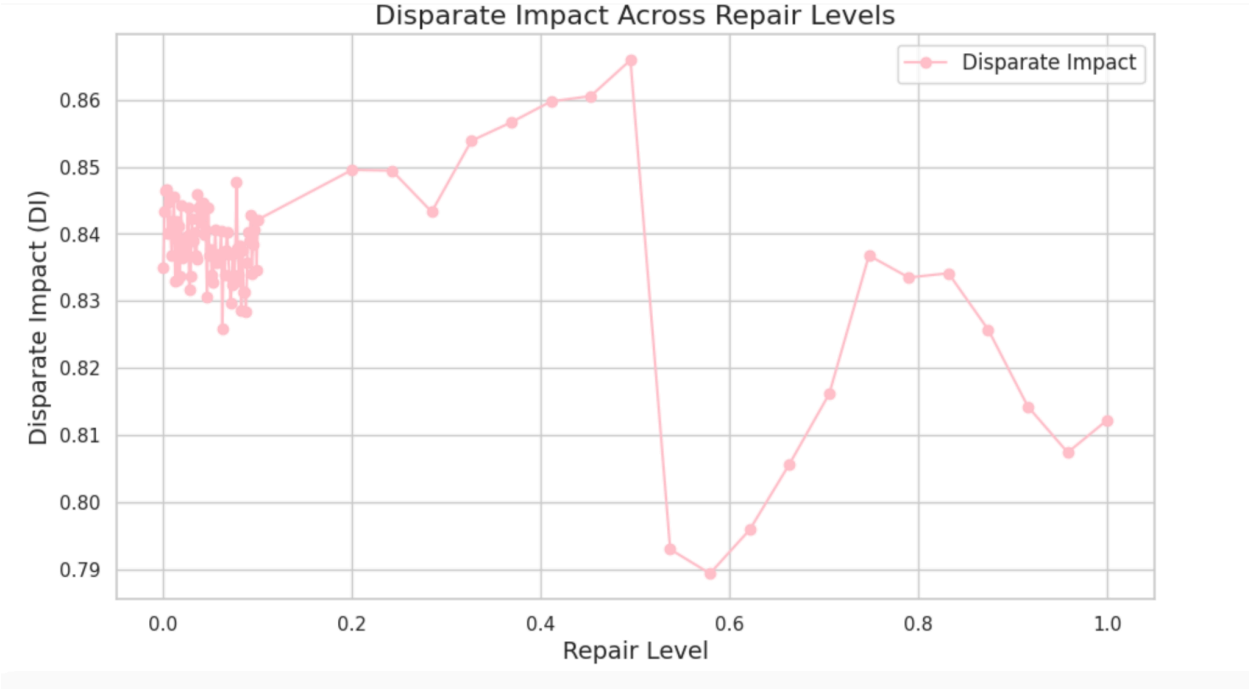

Finer Granular Disparate Impact Results

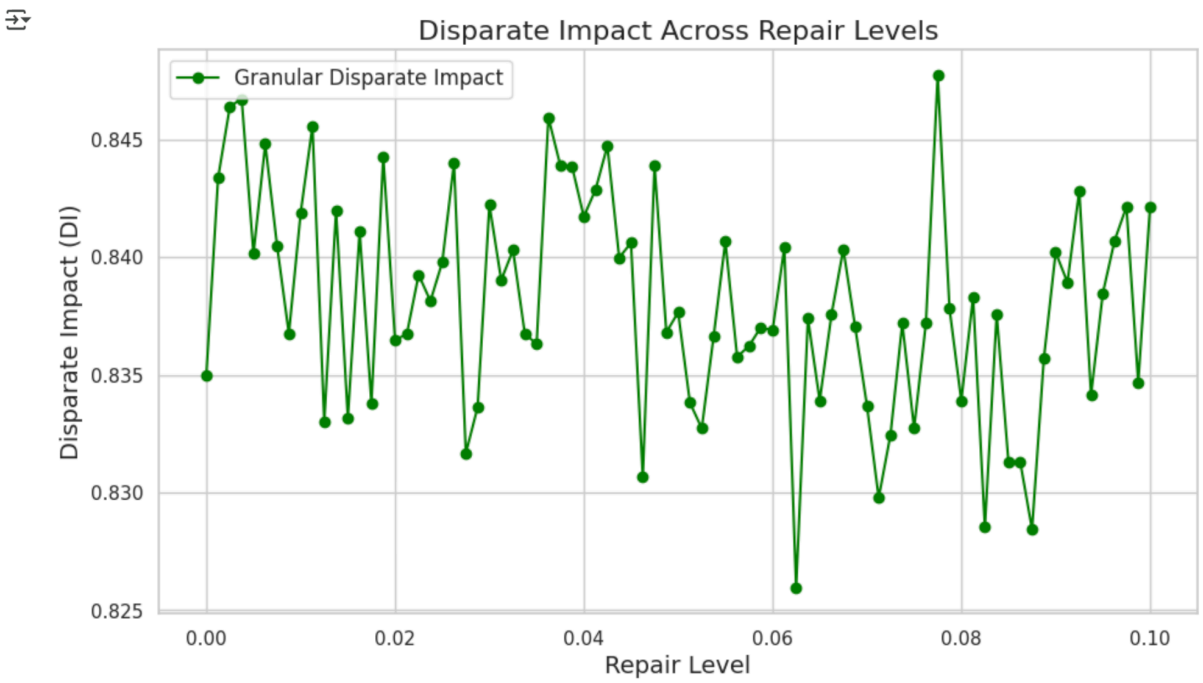

Repair Level: 0.0000, Disparate Impact: 0.8350

Repair Level: 0.0013, Disparate Impact: 0.8434

Repair Level: 0.0025, Disparate Impact: 0.8464

Repair Level: 0.0037, Disparate Impact: 0.8467

Repair Level: 0.0050, Disparate Impact: 0.8401

Repair Level: 0.0063, Disparate Impact: 0.8448

Repair Level: 0.0075, Disparate Impact: 0.8405

Repair Level: 0.0088, Disparate Impact: 0.8367

Repair Level: 0.0100, Disparate Impact: 0.8419

Repair Level: 0.0112, Disparate Impact: 0.8456

Repair Level: 0.0125, Disparate Impact: 0.8330

Repair Level: 0.0138, Disparate Impact: 0.8420

Repair Level: 0.0150, Disparate Impact: 0.8332

Repair Level: 0.0163, Disparate Impact: 0.8411

Repair Level: 0.0175, Disparate Impact: 0.8338

Repair Level: 0.0187, Disparate Impact: 0.8443

Repair Level: 0.0200, Disparate Impact: 0.8365

Repair Level: 0.0213, Disparate Impact: 0.8367

Repair Level: 0.0225, Disparate Impact: 0.8393

Repair Level: 0.0238, Disparate Impact: 0.8381

Repair Level: 0.0250, Disparate Impact: 0.8398

Repair Level: 0.0262, Disparate Impact: 0.8440

Repair Level: 0.0275, Disparate Impact: 0.8317

Repair Level: 0.0288, Disparate Impact: 0.8336

Repair Level: 0.0300, Disparate Impact: 0.8422

Repair Level: 0.0312, Disparate Impact: 0.8390

Repair Level: 0.0325, Disparate Impact: 0.8403

Repair Level: 0.0338, Disparate Impact: 0.8367

Repair Level: 0.0350, Disparate Impact: 0.8363

Repair Level: 0.0362, Disparate Impact: 0.8459

Repair Level: 0.0375, Disparate Impact: 0.8439

Repair Level: 0.0387, Disparate Impact: 0.8439

Repair Level: 0.0400, Disparate Impact: 0.8417

Repair Level: 0.0413, Disparate Impact: 0.8429

Repair Level: 0.0425, Disparate Impact: 0.8447

Repair Level: 0.0438, Disparate Impact: 0.8400

Repair Level: 0.0450, Disparate Impact: 0.8406

Repair Level: 0.0462, Disparate Impact: 0.8307

Repair Level: 0.0475, Disparate Impact: 0.8439

Repair Level: 0.0488, Disparate Impact: 0.8368

Repair Level: 0.0500, Disparate Impact: 0.8377

Repair Level: 0.0513, Disparate Impact: 0.8338

Repair Level: 0.0525, Disparate Impact: 0.8327

Repair Level: 0.0537, Disparate Impact: 0.8366

Repair Level: 0.0550, Disparate Impact: 0.8407

Repair Level: 0.0563, Disparate Impact: 0.8358

Repair Level: 0.0575, Disparate Impact: 0.8362

Repair Level: 0.0588, Disparate Impact: 0.8370

Repair Level: 0.0600, Disparate Impact: 0.8369

Repair Level: 0.0612, Disparate Impact: 0.8404

Repair Level: 0.0625, Disparate Impact: 0.8259

Repair Level: 0.0638, Disparate Impact: 0.8374

Repair Level: 0.0650, Disparate Impact: 0.8339

Repair Level: 0.0663, Disparate Impact: 0.8376

Repair Level: 0.0675, Disparate Impact: 0.8403

Repair Level: 0.0688, Disparate Impact: 0.8370

Repair Level: 0.0700, Disparate Impact: 0.8337

Repair Level: 0.0713, Disparate Impact: 0.8298

Repair Level: 0.0725, Disparate Impact: 0.8325

Repair Level: 0.0737, Disparate Impact: 0.8372

Repair Level: 0.0750, Disparate Impact: 0.8327

Repair Level: 0.0762, Disparate Impact: 0.8372

Repair Level: 0.0775, Disparate Impact: 0.8477

Repair Level: 0.0788, Disparate Impact: 0.8378

Repair Level: 0.0800, Disparate Impact: 0.8339

Repair Level: 0.0813, Disparate Impact: 0.8383

Repair Level: 0.0825, Disparate Impact: 0.8285

Repair Level: 0.0838, Disparate Impact: 0.8376

Repair Level: 0.0850, Disparate Impact: 0.8313

Repair Level: 0.0863, Disparate Impact: 0.8313

Repair Level: 0.0875, Disparate Impact: 0.8284

Repair Level: 0.0887, Disparate Impact: 0.8357

Repair Level: 0.0900, Disparate Impact: 0.8402

Repair Level: 0.0912, Disparate Impact: 0.8389

Repair Level: 0.0925, Disparate Impact: 0.8428

Repair Level: 0.0938, Disparate Impact: 0.8341

Repair Level: 0.0950, Disparate Impact: 0.8385

```
Repair Level: 0.0963, Disparate Impact: 0.8407

Repair Level: 0.0975, Disparate Impact: 0.8421

Repair Level: 0.0988, Disparate Impact: 0.8347

Repair Level: 0.1000, Disparate Impact: 0.8421
```

Result Interpretation of Finer Granular Changes

DI fluctuates significantly, indicating that small changes in fairness repair can lead to inconsistent effects. There is no clear increasing or decreasing trend, and some repair levels improve DI while others reduce it. The DI values stay within a narrow range (0.825 - 0.846), suggesting that fine-grained adjustments may not always yield stable fairness improvement**s**.

1.1b. Code Output for Coarser Granularity

Answer:

The output is organized as below:

- **Coarse Disparate Impact Plots**
- **Coarse Disparate Impact Results**

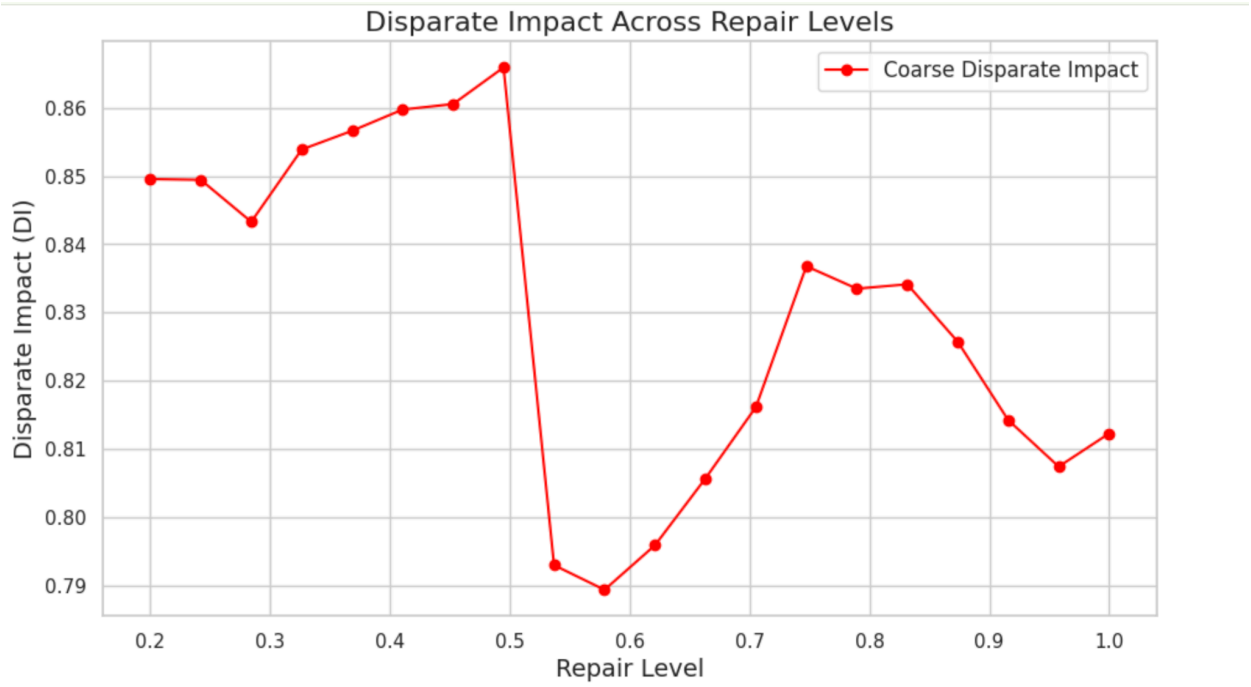

Coarse Disparate Impact Results

```
Repair Level: 0.2000, Disparate Impact: 0.8495
```

```
Repair Level: 0.2421, Disparate Impact: 0.8494
Repair Level: 0.2842, Disparate Impact: 0.8433
Repair Level: 0.3263, Disparate Impact: 0.8539
Repair Level: 0.3684, Disparate Impact: 0.8566
Repair Level: 0.4105, Disparate Impact: 0.8597
Repair Level: 0.4526, Disparate Impact: 0.8605
Repair Level: 0.4947, Disparate Impact: 0.8659
Repair Level: 0.5368, Disparate Impact: 0.7930
Repair Level: 0.5789, Disparate Impact: 0.7894
Repair Level: 0.6211, Disparate Impact: 0.7959
Repair Level: 0.6632, Disparate Impact: 0.8057
Repair Level: 0.7053, Disparate Impact: 0.8161
Repair Level: 0.7474, Disparate Impact: 0.8368
Repair Level: 0.7895, Disparate Impact: 0.8335
Repair Level: 0.8316, Disparate Impact: 0.8341
Repair Level: 0.8737, Disparate Impact: 0.8257
Repair Level: 0.9158, Disparate Impact: 0.8142
Repair Level: 0.9579, Disparate Impact: 0.8074
Repair Level: 1.0000, Disparate Impact: 0.8123
```

Coarse Result Interpretation

Initially, DI remains relatively stable but then increases slightly before experiencing a sharp **drop around the 0.5 repair level**. After this decline, DI gradually recovers but does not fully regain its earlier peak. This suggests that **moderate repair levels improve fairness**, but excessive adjustments (beyond a certain threshold) can have unintended negative effects on fairness. The fluctuations indicate that the **relationship between repair levels and fairness is non-linear**, requiring careful tuning to balance both fairness and performance.

1.2. Code Output for Best Repair Level and Disparate Impact at Best Repair Level

Paste your screenshot here. (Choose the single best out of all repair levels for finer/coarser. This will be the DI closest to 1.0.)

Answer

```
Best Repair Level: 0.4947

Disparate Impact at Best Repair Level: 0.8659
```

1.3. Code Output for Light GBM Performance Metrics and Fairness Metrics

```
Performance Metrics on Test Data (Best DIR Applied):
Accuracy: 0.8103
Precision: 0.6258
Recall: 0.5415
F1-Score: 0.5806
ROC-AUC: 0.7950

Fairness Metrics on Test Data (Best DIR Applied):
Statistical Parity Difference (SPD): -0.1097
Disparate Impact (DI): 0.8659
Equal Opportunity Difference (EOD): -0.0226
Average Odds Difference (AOD): -0.0694
Differential Fairness Bias Amplification (DFBA): 0.1727
```

Part 2: Short-Response Questions

Question:

2.1. How does the Disparate Impact (DI) change with finer granularity? Coarser granularity?

Answer:

Disparate Impact (DI) measures whether an AI model disproportionately favors or disadvantages certain groups. It is calculated as:DI = Selection Rate of Unprivileged Group / Selection Rate of Privileged GroupA DI score closer to 1.0 indicates fairness, while a lower DI suggests bias.

Effect of Finer Granularity on DI:

Finer granularity breaks groups into smaller subcategories for a more detailed assessment.

Example: A bank categorizes borrowers into Male (10% default rate) and Female (15% default rate), resulting in DI = 0.67. If age is added as a factor:

Males 18-30: 12%, 31-50: 9%, 51+: 7%

Females 18-30: 18%, 31-50: 14%, 51+: 11%

Now, DI varies by age, revealing hidden biases. Younger females (DI = 0.67) face more discrimination than older females (DI = 0.78).

Effect of Coarser Granularity on DI:

Coarser granularity merges subgroups, making DI appear more stable but potentially hiding disparities.

Example: If all females are grouped together (15% default rate), age-based bias is masked, and DI remains 0.67 without showing variation.

Conclusion

Finer Granularity → More precise DI but higher fluctuation.

Coarser Granularity → Smoother DI but may hide subgroup biases.

Balanced Approach → Detects bias while minimizing noise.

Question:

2.2. Based on your findings, which model would you recommend (Original, Reweighted, or DIR)? Justify your choice by balancing performance and fairness.

Answer:

The Reweighted Model is the best recommendation because: It achieves the best fairness score (DI closest to 1.0). It still maintains good performance without a major drop in accuracy or precision. It balances both fairness and model effectiveness better than the Original and DIR models.

The results for the all 3 models (Original, Reweighted and DIR) are provided below:

--- Original Model ---

Performance Metrics:

accuracy: 0.8153

precision: 0.6400

recall: 0.5454

f1: 0.5889

roc_auc: 0.8003

Fairness Metrics:

SPD: -0.0862

DI: 0.8972

EOD: 0.0164

AOD: -0.0163

Fairness Score (DI deviation): 0.1028

--- Reweighted Model ---

Performance Metrics:

accuracy: 0.8111

precision: 0.6275

recall: 0.5437

f1: 0.5826

roc_auc: 0.7960

Fairness Metrics:

SPD: -0.0351

DI: 0.9565

EOD: 0.0164

AOD: -0.0163

Fairness Score (DI deviation): 0.0435

--- DIR Model ---

Performance Metrics:

accuracy: 0.8103

precision: 0.6258

recall: 0.5415

f1: 0.5806

roc_auc: 0.7950

Fairness Metrics:

SPD: -0.1097

DI: 0.8659

EOD: -0.0226

AOD: -0.0694

Fairness Score (DI deviation): 0.1341

Recommended Model: Reweighted Model

Question:

2.3. Consider a scenario where you are responsible for deploying a machine learning model that shows a fairness-accuracy trade-off.

The following reflection explores the ethical dilemmas in machine learning, emphasizing the trade-off between fairness and accuracy. Key considerations include organizational values, societal impact, and the specific application context.

Deploying machine learning models that strike a balance between fairness and accuracy requires a principled approach aligned with ethical standards, social responsibility, and business objectives. Organizations must determine their priorities: if efficiency and profitability are the focus, accuracy is paramount for precise predictions; if inclusivity and ethical responsibility take precedence, fairness must be prioritized to mitigate bias and discrimination. For instance, fairness is crucial in hiring models to foster diversity, whereas fraud detection systems prioritize accuracy to minimize financial risk.

Societal values such as justice and accessibility also influence this balance, as AI models can either reinforce or mitigate systemic inequalities. In high-stakes fields like healthcare, finance, or law enforcement, fairness is essential to prevent harm to marginalized groups. A healthcare model might prioritize accuracy for reliable diagnoses, while a loan approval system must carefully balance fairness and accuracy to promote inclusivity without exposing financial institutions to undue risk.

This trade-off presents ethical challenges: focusing too much on fairness may compromise accuracy (e.g., overlooking fraudulent activity), whereas neglecting fairness can perpetuate biases. Transparency and accountability—supported by frameworks like GDPR—help ensure AI decisions remain explainable and ethically sound.

Ultimately, responsible AI development requires ongoing efforts to optimize both fairness and accuracy, with continuous monitoring to adapt to evolving ethical, societal, and business considerations. A well-designed model should not only deliver strong performance but also uphold principles of equity, inclusivity, and justice.